

# ExampleB

March 27, 2024

## 1 Project enda : Example B

If you haven't already, read Example A first, it is not long. Download `example_b.zip` and run this notebook in the correct python environment.

In this example we will go more in depth, with realistic data and more historical data (~4-5 years). This example is divided in 7 parts: 1. Read and prepare data, check for missing values and gaps 2. Visualize data 3. Feature engineering : datetime and calendar features 4. Portfolio forecast & basic prediction 5. Benchmark with simple evaluation 6. Benchmark with Backtesting 7. Make the prediction

We set ourselves in a setup as if we were **exactly on 2020-11-30**. We want to predict the total consumption of customers for the next few days starting 2020-12-01 at a 30min time-step. We have: - our customer contracts until 2020-11-30 included. - historical load data from 2015-01-01 until 2020-11-15 included. There is a ~15 day time-gap between the last moment for which we have an actual load measure and 'today' (2020-11-30). - weather forecast until 2020-12-11 (11 days). - our TSO's network load forecast until 2020-12-7 (7 days).

In here (example B), we will put all our customers in only 1 group and forecast the next 7 days. We will first construct the dataset and the forecast input data and test it with a basic linear regressor. We will then try various algorithms and compare them. Finally we will give an example of backtesting on the data.

```
[1]: import enda
import pandas as pd
import os
```

```
[2]: enda.__file__
```

```
[2]: '/Users/clement.jeannesson/Jobs/enda/enda/__init__.py'
```

### 1.1 1. Read and prepare data, check for missing values and gaps

```
[3]: # Replace this with the path to your example_b directory.
# You should have ExampleB.ipynb opened in jupyter, so you can run each step
DIR = '.'
```

```
[4]: # Get the 30min time-step data just like in Example A
# (columns are a bit different and there is more data)
```

```

# Here we consider all customers in one big group.
def read_data():
    contracts = enda.Contracts.read_contracts_from_file(os.path.join(DIR,
↪ "contracts.csv"))
    contracts["contracts_count"] = 1
    portfolio_by_day = enda.Contracts.compute_portfolio_by_day(
        contracts,
        columns_to_sum = ["contracts_count", "kva"],
        date_start_col="date_start",
        date_end_exclusive_col="date_end_exclusive",
    )
    portfolio = enda.Resample.upsample_and_interpolate(
        portfolio_by_day,
        freq='30min',
        tz_info='Europe/Paris',
        forward_fill=True
    )

    historic_load_measured = pd.read_csv(os.path.join(DIR,
↪ "historic_load_measured.csv"))
    weather_and_tso_forecasts = pd.read_csv(os.path.join(DIR,
↪ "weather_and_tso_forecasts.csv"))
    # correctly format 'time' as a pandas.DatetimeIndex of dtype: datetime[ns,
↪ tzinfo]
    for df in [historic_load_measured, weather_and_tso_forecasts]:
        df['time'] = pd.to_datetime(df['time'])
        df['time'] = enda.TimezoneUtils.
↪ convert_dtype_from_object_to_tz_aware(df['time'], tz_info = 'Europe/Paris')
        df.set_index('time', inplace=True)

    # keep only where both loads are known
    historic_load_measured = historic_load_measured.dropna()
    historic_load_measured["load_kw"] =
↪ historic_load_measured["smart_metered_kw"] + historic_load_measured["slp_kw"]
    # keep only the full load
    historic_load_measured = historic_load_measured[["load_kw"]]

    return contracts, portfolio, historic_load_measured,
↪ weather_and_tso_forecasts

```

```

[5]: contracts, portfolio, historic_load_measured, weather_and_tso_forecasts =
↪ read_data()

```

```

[6]: contracts

```

```
[6]:
```

	date_start	date_end_exclusive	kva	meter_reading_type	contracts_count
0	2006-08-09	NaT	12.0	PROFILE	1
1	2006-09-01	2006-11-23	6.0	PROFILE	1
2	2006-09-01	2007-11-01	3.0	PROFILE	1
3	2006-09-01	2007-12-19	12.0	PROFILE	1
4	2006-09-01	2008-06-28	12.0	PROFILE	1
...	...	...	...	...	...
162598	2020-11-30	NaT	6.0	PROFILE	1
162599	2020-11-30	NaT	6.0	PROFILE	1
162600	2020-11-30	NaT	6.0	PROFILE	1
162601	2020-11-30	NaT	6.0	PROFILE	1
162602	2020-11-30	NaT	6.0	PROFILE	1

[162603 rows x 5 columns]

```
[7]: portfolio
```

```
[7]:
```

	contracts_count	kva
date		
2006-08-09 00:00:00+02:00	1.0	12.0
2006-08-09 00:30:00+02:00	1.0	12.0
2006-08-09 01:00:00+02:00	1.0	12.0
2006-08-09 01:30:00+02:00	1.0	12.0
2006-08-09 02:00:00+02:00	1.0	12.0
...	...	...
2020-11-30 21:30:00+01:00	96134.0	820005.7
2020-11-30 22:00:00+01:00	96134.0	820005.7
2020-11-30 22:30:00+01:00	96134.0	820005.7
2020-11-30 23:00:00+01:00	96134.0	820005.7
2020-11-30 23:30:00+01:00	96134.0	820005.7

[250946 rows x 2 columns]

```
[8]: historic_load_measured
```

```
[8]:
```

	load_kw
time	
2015-01-01 00:00:00+01:00	2490.925806
2015-01-01 00:30:00+01:00	2412.623113
2015-01-01 01:00:00+01:00	2365.611276
2015-01-01 01:30:00+01:00	2336.141065
2015-01-01 02:00:00+01:00	2300.935642
...	...
2020-11-15 21:30:00+01:00	7657.293444
2020-11-15 22:00:00+01:00	7317.540759
2020-11-15 22:30:00+01:00	7580.051439
2020-11-15 23:00:00+01:00	7496.273993

2020-11-15 23:30:00+01:00 7376.005701

[97198 rows x 1 columns]

```
[9]: # t_weighted is the average french temperature weighted by population density
# t_smooth is a smoothing computed over t_weighted to take into account
↳ building calorific inertia
# (t_smooth is computed out of enda here)

# some tso_forecast_load_mw is missing at the end (we don't show it here)
weather_and_tso_forecasts.dropna(subset=["tso_forecast_load_mw"])
```

```
[9]:
```

	tso_forecast_load_mw	t_weighted	t_smooth
time			
2015-01-01 00:00:00+01:00	72900.0	-0.41	1.17
2015-01-01 00:30:00+01:00	71600.0	-0.48	1.17
2015-01-01 01:00:00+01:00	69900.0	-0.55	1.15
2015-01-01 01:30:00+01:00	70600.0	-0.66	1.14
2015-01-01 02:00:00+01:00	70500.0	-0.78	1.11
...	...	...	...
2020-12-07 21:30:00+01:00	68400.0	4.20	4.13
2020-12-07 22:00:00+01:00	66900.0	4.12	4.10
2020-12-07 22:30:00+01:00	67600.0	4.03	4.08
2020-12-07 23:00:00+01:00	70200.0	3.94	4.07
2020-12-07 23:30:00+01:00	69600.0	3.94	4.07

[104064 rows x 3 columns]

```
[10]: # lets create the train set with historical data
historic = pd.merge(
    portfolio,
    historic_load_measured, # here we select only the load of the desired group
    how='inner', left_index=True, right_index=True
)

historic = pd.merge(
    historic,
    weather_and_tso_forecasts,
    how='inner', left_index=True, right_index=True
)
```

```
[11]: historic
```

```
[11]:
```

	contracts_count	kva	load_kw	\
2015-01-01 00:00:00+01:00	21261.000000	167416.4000	2490.925806	
2015-01-01 00:30:00+01:00	21261.020833	167417.4000	2412.623113	
2015-01-01 01:00:00+01:00	21261.041667	167418.4000	2365.611276	

2015-01-01 01:30:00+01:00	21261.062500	167419.4000	2336.141065
2015-01-01 02:00:00+01:00	21261.083333	167420.4000	2300.935642
...	...	...	...
2020-11-15 21:30:00+01:00	95509.041667	813616.3625	7657.293444
2020-11-15 22:00:00+01:00	95509.833333	813623.0500	7317.540759
2020-11-15 22:30:00+01:00	95510.625000	813629.7375	7580.051439
2020-11-15 23:00:00+01:00	95511.416667	813636.4250	7496.273993
2020-11-15 23:30:00+01:00	95512.208333	813643.1125	7376.005701

	tso_forecast_load_mw	t_weighted	t_smooth
2015-01-01 00:00:00+01:00	72900.0	-0.41	1.17
2015-01-01 00:30:00+01:00	71600.0	-0.48	1.17
2015-01-01 01:00:00+01:00	69900.0	-0.55	1.15
2015-01-01 01:30:00+01:00	70600.0	-0.66	1.14
2015-01-01 02:00:00+01:00	70500.0	-0.78	1.11
...	...	...	...
2020-11-15 21:30:00+01:00	46200.0	12.05	12.01
2020-11-15 22:00:00+01:00	45200.0	11.92	11.97
2020-11-15 22:30:00+01:00	46400.0	11.84	11.96
2020-11-15 23:00:00+01:00	48600.0	11.75	11.94
2020-11-15 23:30:00+01:00	49400.0	11.64	11.92

[97198 rows x 6 columns]

```
[12]: # check that there is no NaN value
historic.isna().sum()
```

```
[12]: contracts_count      0
kva                        0
load_kw                   0
tso_forecast_load_mw      0
t_weighted                0
t_smooth                  0
dtype: int64
```

```
[13]: # note that the type of the index is precise
historic.index.dtype, type(historic.index)
```

```
[13]: (datetime64[ns, Europe/Paris], pandas.core.indexes.datetimes.DatetimeIndex)
```

```
[14]: # check missing data in the timeseries (based on the time index only)
missing_periods= enda.TimeSeries.find_missing_periods(
    dti=historic.index,
    expected_freq = '30min',
    expected_start_datetime = pd.to_datetime('2015-01-01 00:00:00+01:00').
    ↳astimezone('Europe/Paris'),
```

```

        expected_excl_end_datetime = pd.to_datetime('2020-12-01 00:00:00+01:00').
        ↪astimezone('Europe/Paris'),
    )
    for missing_period in missing_periods:
        print("Missing data from {} to {}".format(missing_period[0], ↪
        ↪missing_period[1]))

    duplicates, extra_points = enda.TimeSeries.find_duplicates_and_extra_points(
        dti=historic.index,
        expected_freq = '30min'
    )

    if len(duplicates) > 0 :
        print("Extra points found: {}".format(duplicates))

    if len(extra_points) > 0 :
        print("Extra points found: {}".format(extra_points))

```

Missing data from 2015-09-01 00:00:00+02:00 to 2015-11-30 23:30:00+01:00.

Missing data from 2018-06-01 00:00:00+02:00 to 2018-06-30 23:30:00+02:00.

Missing data from 2020-11-16 00:00:00+01:00 to 2020-11-30 23:30:00+01:00.

We expect some missing data.

```

[15]: # Zoom on a daylight savings time change to double-check that it was handled ↪
        ↪correctly
    historic[(historic.index >= '2019-10-27 01:00:00+02:00') & (historic.index < ↪
        ↪'2019-10-27 03:30:00+01:00')]

```

```

[15]:

```

	contracts_count	kva	load_kw	\
2019-10-27 01:00:00+02:00	84133.24	716839.24	5179.955556	
2019-10-27 01:30:00+02:00	84134.36	716850.66	5087.111111	
2019-10-27 02:00:00+02:00	84135.48	716862.08	4898.400000	
2019-10-27 02:30:00+02:00	84136.60	716873.50	4616.533333	
2019-10-27 02:00:00+01:00	84137.72	716884.92	4259.822222	
2019-10-27 02:30:00+01:00	84138.84	716896.34	4208.888889	
2019-10-27 03:00:00+01:00	84139.96	716907.76	4137.955556	

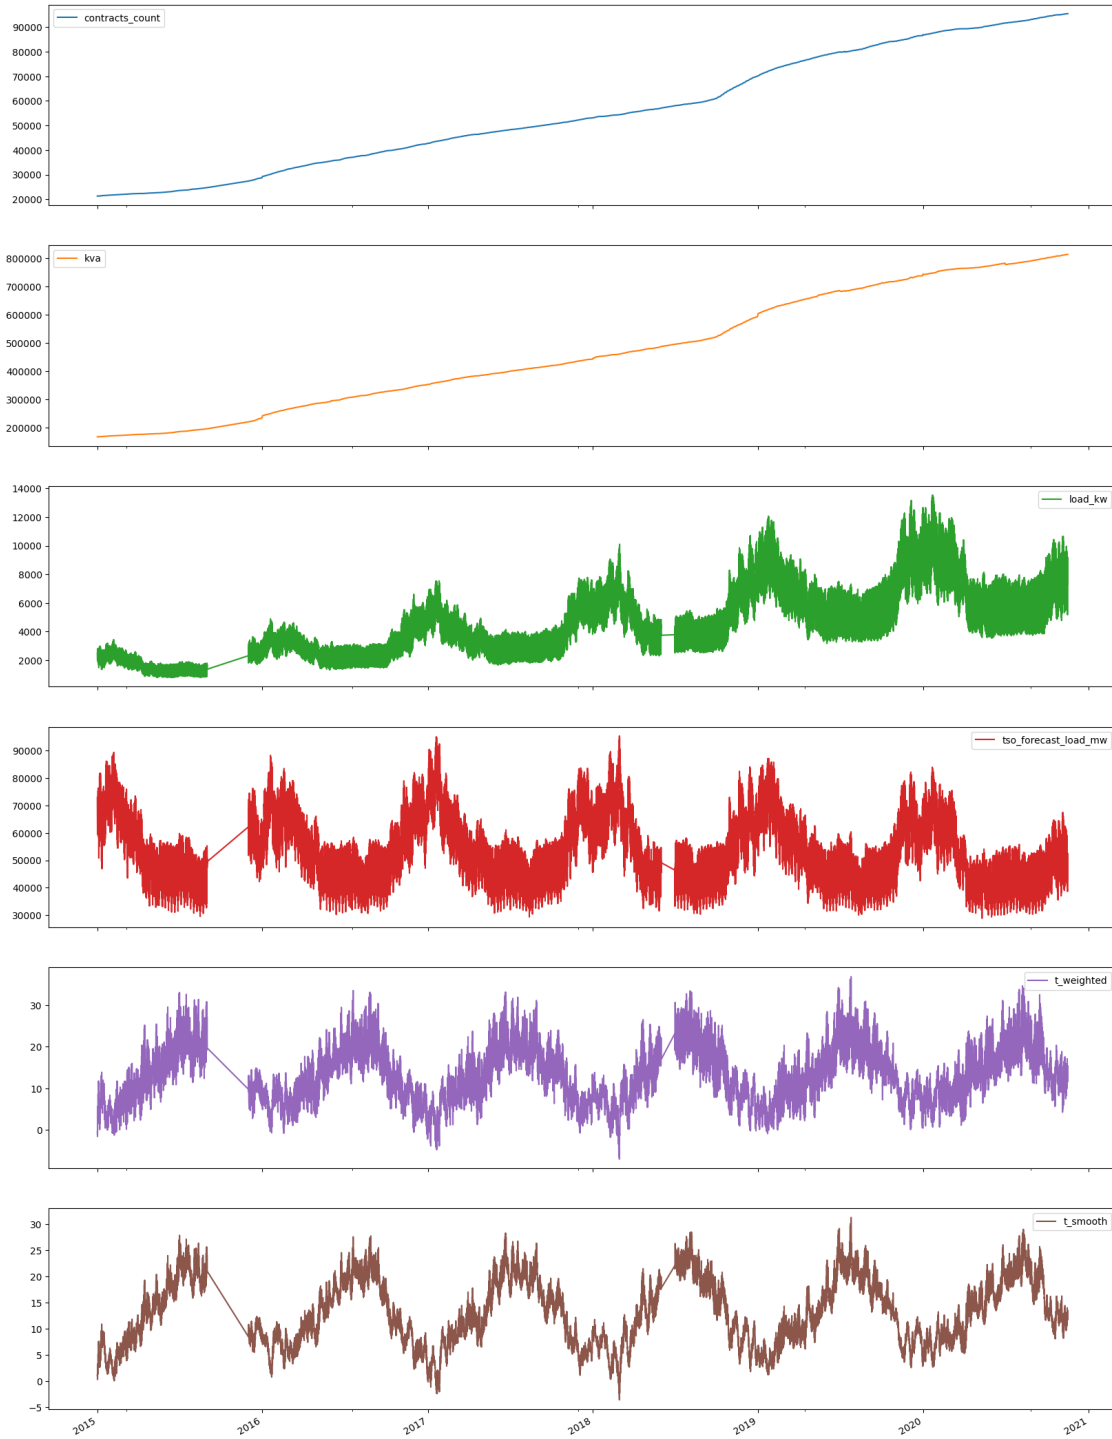
	tso_forecast_load_mw	t_weighted	t_smooth
2019-10-27 01:00:00+02:00	41300.0	13.65	13.49
2019-10-27 01:30:00+02:00	40700.0	13.52	13.47
2019-10-27 02:00:00+02:00	36700.0	13.40	13.46
2019-10-27 02:30:00+02:00	36700.0	13.26	13.44
2019-10-27 02:00:00+01:00	36700.0	13.12	13.42
2019-10-27 02:30:00+01:00	36700.0	12.91	13.39
2019-10-27 03:00:00+01:00	36700.0	12.70	13.37

## 1.2 2. Visualize data

In order to visualise using pandas, we use the matplotlib backend.

```
[16]: # Show full data set  
historic.plot(figsize=(20, 30), subplots=True)
```


```
[16]: array([<Axes: >, <Axes: >, <Axes: >, <Axes: >, <Axes: >, <Axes: >],  
          dtype=object)
```



```
[17]: # Show recent data
historic[-1000:].plot(figsize=(20, 30), subplots=True)
```



```
[17]: array([<Axes: >, <Axes: >, <Axes: >, <Axes: >, <Axes: >, <Axes: >],  
          dtype=object)
```



Don't hesitate to add your own visualisations!

### 1.3 3. Feature engineering

Before we train, we will add some features based on the `datetime`, and some calendar features related to national holidays or school holidays.

We use some packages for the holidays, which are used in `enda.feature_engineering.calendar`:

```
[18]: def featurize_datetime(df: pd.DataFrame) -> pd.DataFrame:
      """
      Featurize the input dataframe with date/datetime-oriented features
      """

      # make a copy
      df = df.copy()

      df = enda.DatetimeFeature.split_datetime(
          df, split_list = ['minuteofday', 'dayofweek']
      )
      df = enda.DatetimeFeature.encode_cyclic_datetime_index(
          df, split_list = ['minuteofday', 'dayofweek', 'dayofyear']
      )

      # min and max years
      min_year = min(df.index).year
      max_year = max(df.index).year

      # add features about national holidays and school holidays (French holidays
      ↪ here)
      special_days = enda.Calendar.feature_special_days(country='FR', years_list
      ↪ = [min_year, max_year+1])
      df = pd.merge(
          df,
          special_days,
          how='left',
          left_index=True,
          right_index=True
      )

      return df
```

```
[19]: # feature the historic dataframe. It makes it a train_set.
      full_train_set = featurize_datetime(historic)
```

```
[20]: full_train_set
```

[20]:

	contracts_count	kva	load_kw	\
2015-01-01 00:00:00+01:00	21261.000000	167416.4000	2490.925806	
2015-01-01 00:30:00+01:00	21261.020833	167417.4000	2412.623113	
2015-01-01 01:00:00+01:00	21261.041667	167418.4000	2365.611276	
2015-01-01 01:30:00+01:00	21261.062500	167419.4000	2336.141065	
2015-01-01 02:00:00+01:00	21261.083333	167420.4000	2300.935642	
...	...	...	...	
2020-11-15 21:30:00+01:00	95509.041667	813616.3625	7657.293444	
2020-11-15 22:00:00+01:00	95509.833333	813623.0500	7317.540759	
2020-11-15 22:30:00+01:00	95510.625000	813629.7375	7580.051439	
2020-11-15 23:00:00+01:00	95511.416667	813636.4250	7496.273993	
2020-11-15 23:30:00+01:00	95512.208333	813643.1125	7376.005701	

	tso_forecast_load_mw	t_weighted	t_smooth	\
2015-01-01 00:00:00+01:00	72900.0	-0.41	1.17	
2015-01-01 00:30:00+01:00	71600.0	-0.48	1.17	
2015-01-01 01:00:00+01:00	69900.0	-0.55	1.15	
2015-01-01 01:30:00+01:00	70600.0	-0.66	1.14	
2015-01-01 02:00:00+01:00	70500.0	-0.78	1.11	
...	...	...	...	
2020-11-15 21:30:00+01:00	46200.0	12.05	12.01	
2020-11-15 22:00:00+01:00	45200.0	11.92	11.97	
2020-11-15 22:30:00+01:00	46400.0	11.84	11.96	
2020-11-15 23:00:00+01:00	48600.0	11.75	11.94	
2020-11-15 23:30:00+01:00	49400.0	11.64	11.92	

	minuteofday	dayofweek	minuteofday_cos	\
2015-01-01 00:00:00+01:00	0	3	1.000000	
2015-01-01 00:30:00+01:00	30	3	0.991445	
2015-01-01 01:00:00+01:00	60	3	0.965926	
2015-01-01 01:30:00+01:00	90	3	0.923880	
2015-01-01 02:00:00+01:00	120	3	0.866025	
...	...	...	...	
2020-11-15 21:30:00+01:00	1290	6	0.793353	
2020-11-15 22:00:00+01:00	1320	6	0.866025	
2020-11-15 22:30:00+01:00	1350	6	0.923880	
2020-11-15 23:00:00+01:00	1380	6	0.965926	
2020-11-15 23:30:00+01:00	1410	6	0.991445	

	minuteofday_sin	dayofweek_cos	dayofweek_sin	\
2015-01-01 00:00:00+01:00	0.000000	-0.900969	0.433884	
2015-01-01 00:30:00+01:00	0.130526	-0.900969	0.433884	
2015-01-01 01:00:00+01:00	0.258819	-0.900969	0.433884	
2015-01-01 01:30:00+01:00	0.382683	-0.900969	0.433884	
2015-01-01 02:00:00+01:00	0.500000	-0.900969	0.433884	
...	...	...	...	
2020-11-15 21:30:00+01:00	-0.608761	0.623490	-0.781831	

2020-11-15 22:00:00+01:00	-0.500000	0.623490	-0.781831
2020-11-15 22:30:00+01:00	-0.382683	0.623490	-0.781831
2020-11-15 23:00:00+01:00	-0.258819	0.623490	-0.781831
2020-11-15 23:30:00+01:00	-0.130526	0.623490	-0.781831

	dayofyear_cos	dayofyear_sin	lockdown \
2015-01-01 00:00:00+01:00	1.000000	0.000000	0.0
2015-01-01 00:30:00+01:00	1.000000	0.000000	0.0
2015-01-01 01:00:00+01:00	1.000000	0.000000	0.0
2015-01-01 01:30:00+01:00	1.000000	0.000000	0.0
2015-01-01 02:00:00+01:00	1.000000	0.000000	0.0
...	...	...	...
2020-11-15 21:30:00+01:00	0.691771	-0.722117	0.0
2020-11-15 22:00:00+01:00	0.691771	-0.722117	0.0
2020-11-15 22:30:00+01:00	0.691771	-0.722117	0.0
2020-11-15 23:00:00+01:00	0.691771	-0.722117	0.0
2020-11-15 23:30:00+01:00	0.691771	-0.722117	0.0

	public_holiday	nb_school_areas_off \
2015-01-01 00:00:00+01:00	1.0	3.0
2015-01-01 00:30:00+01:00	1.0	3.0
2015-01-01 01:00:00+01:00	1.0	3.0
2015-01-01 01:30:00+01:00	1.0	3.0
2015-01-01 02:00:00+01:00	1.0	3.0
...	...	...
2020-11-15 21:30:00+01:00	0.0	0.0
2020-11-15 22:00:00+01:00	0.0	0.0
2020-11-15 22:30:00+01:00	0.0	0.0
2020-11-15 23:00:00+01:00	0.0	0.0
2020-11-15 23:30:00+01:00	0.0	0.0

	extra_long_weekend
2015-01-01 00:00:00+01:00	0.0
2015-01-01 00:30:00+01:00	0.0
2015-01-01 01:00:00+01:00	0.0
2015-01-01 01:30:00+01:00	0.0
2015-01-01 02:00:00+01:00	0.0
...	...
2020-11-15 21:30:00+01:00	0.0
2020-11-15 22:00:00+01:00	0.0
2020-11-15 22:30:00+01:00	0.0
2020-11-15 23:00:00+01:00	0.0
2020-11-15 23:30:00+01:00	0.0

[97198 rows x 18 columns]

```
[21]: # train a basic scikit-learn LinearRegression
from enda.ml_backends.sklearn_estimator import EndaSklearnEstimator
from sklearn.linear_model import LinearRegression

lin_reg = EndaSklearnEstimator(LinearRegression())
lin_reg.train(full_train_set, target_col='load_kw')
```

## 1.4 4. Portfolio forecast & basic prediction

We need an estimate of our portfolio in the next few days, the tso\_load and weather forecasts.

In order to get our portfolio in the next few days, here we will just consider the latest trends in our portfolio.

In another setup, you might want to connect to your sales software or ERP and take into account contracts that will end or start soon.

We will use `enda.Contracts.forecast_portfolio_linear` (which requires the `sklearn` package).

```
[22]: # we will forecast the portfolio using a linear method
forecast_portfolio = enda.Contracts.forecast_portfolio_linear(
    portfolio_df=portfolio[portfolio.index >= "2020-11-01 00:00:00+02:00"], # #
    ↪only use recent portfolio trend to forecast the next few days
    start_forecast_date=pd.to_datetime("2020-12-01 00:00:00+01:00").
    ↪tz_convert("Europe/Paris"),
    end_forecast_date_exclusive=pd.to_datetime("2020-12-08 00:00:00+01:00").
    ↪tz_convert("Europe/Paris"),
    freq='30min',
    tzinfo='Europe/Paris'
)
forecast_portfolio
```

```
[22]:
```

	contracts_count	kva
date		
2020-12-01 00:00:00+01:00	96045.999741	819322.586540
2020-12-01 00:30:00+01:00	96046.649342	819329.296753
2020-12-01 01:00:00+01:00	96047.298944	819336.006967
2020-12-01 01:30:00+01:00	96047.948545	819342.717181
2020-12-01 02:00:00+01:00	96048.598146	819349.427394
...	...	...
2020-12-07 21:30:00+01:00	96261.017767	821543.667271
2020-12-07 22:00:00+01:00	96261.667368	821550.377484
2020-12-07 22:30:00+01:00	96262.316969	821557.087698
2020-12-07 23:00:00+01:00	96262.966571	821563.797912
2020-12-07 23:30:00+01:00	96263.616172	821570.508125

[336 rows x 2 columns]

```
[23]: # add weather_and_tso_forecasts
forecast_input_data = pd.merge(
    forecast_portfolio,
    weather_and_tso_forecasts.dropna(subset=["tso_forecast_load_mw"]), #
    ↪forecast only where tso is not null for now
    how='inner', left_index=True, right_index=True
)
# add feature engineering
forecast_input_data = featurize_datetime(forecast_input_data)
forecast_input_data
```

```
[23]:
```

	contracts_count	kva \
2020-12-01 00:00:00+01:00	96045.999741	819322.586540
2020-12-01 00:30:00+01:00	96046.649342	819329.296753
2020-12-01 01:00:00+01:00	96047.298944	819336.006967
2020-12-01 01:30:00+01:00	96047.948545	819342.717181
2020-12-01 02:00:00+01:00	96048.598146	819349.427394
...	...	...
2020-12-07 21:30:00+01:00	96261.017767	821543.667271
2020-12-07 22:00:00+01:00	96261.667368	821550.377484
2020-12-07 22:30:00+01:00	96262.316969	821557.087698
2020-12-07 23:00:00+01:00	96262.966571	821563.797912
2020-12-07 23:30:00+01:00	96263.616172	821570.508125

	tso_forecast_load_mw	t_weighted	t_smooth \
2020-12-01 00:00:00+01:00	66100.0	4.69	5.08
2020-12-01 00:30:00+01:00	64200.0	4.82	5.10
2020-12-01 01:00:00+01:00	61900.0	4.96	5.12
2020-12-01 01:30:00+01:00	62800.0	5.04	5.13
2020-12-01 02:00:00+01:00	62300.0	5.13	5.14
...	...	...	...
2020-12-07 21:30:00+01:00	68400.0	4.20	4.13
2020-12-07 22:00:00+01:00	66900.0	4.12	4.10
2020-12-07 22:30:00+01:00	67600.0	4.03	4.08
2020-12-07 23:00:00+01:00	70200.0	3.94	4.07
2020-12-07 23:30:00+01:00	69600.0	3.94	4.07

	minuteofday	dayofweek	minuteofday_cos \
2020-12-01 00:00:00+01:00	0	1	1.000000
2020-12-01 00:30:00+01:00	30	1	0.991445
2020-12-01 01:00:00+01:00	60	1	0.965926
2020-12-01 01:30:00+01:00	90	1	0.923880
2020-12-01 02:00:00+01:00	120	1	0.866025
...	...	...	...
2020-12-07 21:30:00+01:00	1290	0	0.793353
2020-12-07 22:00:00+01:00	1320	0	0.866025
2020-12-07 22:30:00+01:00	1350	0	0.923880

2020-12-07 23:00:00+01:00	1380	0	0.965926
2020-12-07 23:30:00+01:00	1410	0	0.991445

	minuteofday_sin	dayofweek_cos	dayofweek_sin \
2020-12-01 00:00:00+01:00	0.000000	0.62349	0.781831
2020-12-01 00:30:00+01:00	0.130526	0.62349	0.781831
2020-12-01 01:00:00+01:00	0.258819	0.62349	0.781831
2020-12-01 01:30:00+01:00	0.382683	0.62349	0.781831
2020-12-01 02:00:00+01:00	0.500000	0.62349	0.781831
...	...	...	...
2020-12-07 21:30:00+01:00	-0.608761	1.00000	0.000000
2020-12-07 22:00:00+01:00	-0.500000	1.00000	0.000000
2020-12-07 22:30:00+01:00	-0.382683	1.00000	0.000000
2020-12-07 23:00:00+01:00	-0.258819	1.00000	0.000000
2020-12-07 23:30:00+01:00	-0.130526	1.00000	0.000000

	dayofyear_cos	dayofyear_sin	lockdown \
2020-12-01 00:00:00+01:00	0.861702	-0.507415	0.0
2020-12-01 00:30:00+01:00	0.861702	-0.507415	0.0
2020-12-01 01:00:00+01:00	0.861702	-0.507415	0.0
2020-12-01 01:30:00+01:00	0.861702	-0.507415	0.0
2020-12-01 02:00:00+01:00	0.861702	-0.507415	0.0
...	...	...	...
2020-12-07 21:30:00+01:00	0.909308	-0.416125	0.0
2020-12-07 22:00:00+01:00	0.909308	-0.416125	0.0
2020-12-07 22:30:00+01:00	0.909308	-0.416125	0.0
2020-12-07 23:00:00+01:00	0.909308	-0.416125	0.0
2020-12-07 23:30:00+01:00	0.909308	-0.416125	0.0

	public_holiday	nb_school_areas_off \
2020-12-01 00:00:00+01:00	0.0	0.0
2020-12-01 00:30:00+01:00	0.0	0.0
2020-12-01 01:00:00+01:00	0.0	0.0
2020-12-01 01:30:00+01:00	0.0	0.0
2020-12-01 02:00:00+01:00	0.0	0.0
...	...	...
2020-12-07 21:30:00+01:00	0.0	0.0
2020-12-07 22:00:00+01:00	0.0	0.0
2020-12-07 22:30:00+01:00	0.0	0.0
2020-12-07 23:00:00+01:00	0.0	0.0
2020-12-07 23:30:00+01:00	0.0	0.0

	extra_long_weekend
2020-12-01 00:00:00+01:00	0.0
2020-12-01 00:30:00+01:00	0.0
2020-12-01 01:00:00+01:00	0.0
2020-12-01 01:30:00+01:00	0.0



```

2020-12-01 02:00:00+01:00      0.0
...
2020-12-07 21:30:00+01:00      0.0
2020-12-07 22:00:00+01:00      0.0
2020-12-07 22:30:00+01:00      0.0
2020-12-07 23:00:00+01:00      0.0
2020-12-07 23:30:00+01:00      0.0

```



[336 rows x 17 columns]

```

[24]: # show recent portfolio and forecast
for c in ["contracts_count", "kva"]:
    to_plot = pd.merge(
        portfolio[(portfolio.index >= '2020-10-01')][c].to_frame(c+"_historic"),
        forecast_input_data[c].to_frame(c+"_forecast"),
        how='outer', left_index=True, right_index=True
    )

    to_plot.plot(figsize=(16, 4))

```



```

[25]: # do the prediction
lin_reg_prediction = lin_reg.predict(forecast_input_data, target_col="load_kw")

```

```
[26]: lin_reg_prediction
```


```
[26]:
```

	load_kw
2020-12-01 00:00:00+01:00	8990.777945
2020-12-01 00:30:00+01:00	8786.112034
2020-12-01 01:00:00+01:00	8548.345484
2020-12-01 01:30:00+01:00	8630.046223
2020-12-01 02:00:00+01:00	8581.896054
...	...
2020-12-07 21:30:00+01:00	9891.252513
2020-12-07 22:00:00+01:00	9709.005148
2020-12-07 22:30:00+01:00	9744.242062
2020-12-07 23:00:00+01:00	9968.246215
2020-12-07 23:30:00+01:00	9882.876502

[336 rows x 1 columns]

```
[27]: # visualize recent load along with our forecast.  
# remember we don't have recent actual load so there is a time-gap.  
to_plot = pd.merge(  
    historic["load_kw"][-4000:].to_frame("historic_kw"),  
    lin_reg_prediction.rename(columns={"load_kw": "forecast_kw"}),  
    how='outer', left_index=True, right_index=True  
)  
to_plot.plot(ylim=0, figsize=(16, 4))
```

```
[27]: <Axes: >
```



## 1.5 5. Benchmark with simple evaluation

The previous forecast based on linear regression is very limited. Let's try and use a better algorithm !

We will define some algorithms using `scikit-learn` as a machine learning backend and others using `h2o`.

For that we need the `h2o` package:

```
[28]: # here we do a benchmark, we want to compare with actual data,
# lets say from 2020-11-01 to 2020-11-15
benchmark_train = full_train_set[full_train_set.index < '2020-11-01']
benchmark_test = full_train_set[full_train_set.index >= '2020-11-01']

# save the actual_load in a 'benchmark' dataframe,
# we will add the predictions of each algo to 'benchmark'
benchmark = benchmark_test["load_kw"].to_frame("actual_load_kw")

benchmark_test = benchmark_test.drop(columns=["load_kw"])
benchmark
```

```
[28]:
```

	actual_load_kw
2020-11-01 00:00:00+01:00	6817.332090
2020-11-01 00:30:00+01:00	6326.667322
2020-11-01 01:00:00+01:00	6172.223671
2020-11-01 01:30:00+01:00	6050.575318
2020-11-01 02:00:00+01:00	5898.881230
...	...
2020-11-15 21:30:00+01:00	7657.293444
2020-11-15 22:00:00+01:00	7317.540759
2020-11-15 22:30:00+01:00	7580.051439
2020-11-15 23:00:00+01:00	7496.273993
2020-11-15 23:30:00+01:00	7376.005701

[720 rows x 1 columns]

```
[29]: # some parts give ConvergenceWarnings here and we'll ignore them.
import warnings
warnings.filterwarnings('ignore')
```

```
[30]: # use the same method as before to predict a portfolio for 2020-11-01 ->
↳2020-11-15
benchmark_test_portfolio = forecast_portfolio = enda.Contracts.
↳forecast_portfolio_linear(
    portfolio_df=portfolio[(portfolio.index >= '2020-10-01') & (portfolio.index_
↳< '2020-11-01')],
    start_forecast_date=pd.to_datetime("2020-11-01 00:00:00+01:00").
↳tz_convert("Europe/Paris"),
    end_forecast_date_exclusive=pd.to_datetime("2020-11-16 00:00:00+01:00").
↳tz_convert("Europe/Paris"),
    freq='30min',
    tzinfo='Europe/Paris'
)

benchmark_test['kva'] = benchmark_test_portfolio['kva']
benchmark_test['contracts_count'] = benchmark_test_portfolio['contracts_count']
```

```
benchmark_test
```

[30]:

	contracts_count	kva \
2020-11-01 00:00:00+01:00	95205.814480	809817.741508
2020-11-01 00:30:00+01:00	95206.326320	809823.316604
2020-11-01 01:00:00+01:00	95206.838160	809828.891701
2020-11-01 01:30:00+01:00	95207.350000	809834.466797
2020-11-01 02:00:00+01:00	95207.861839	809840.041893
...	...	...
2020-11-15 21:30:00+01:00	95571.779928	813803.935204
2020-11-15 22:00:00+01:00	95572.291767	813809.510300
2020-11-15 22:30:00+01:00	95572.803607	813815.085396
2020-11-15 23:00:00+01:00	95573.315447	813820.660492
2020-11-15 23:30:00+01:00	95573.827287	813826.235588

	tso_forecast_load_mw	t_weighted	t_smooth \
2020-11-01 00:00:00+01:00	47900.0	12.67	12.37
2020-11-01 00:30:00+01:00	45800.0	12.68	12.37
2020-11-01 01:00:00+01:00	43700.0	12.70	12.37
2020-11-01 01:30:00+01:00	43900.0	12.66	12.37
2020-11-01 02:00:00+01:00	43200.0	12.63	12.36
...	...	...	...
2020-11-15 21:30:00+01:00	46200.0	12.05	12.01
2020-11-15 22:00:00+01:00	45200.0	11.92	11.97
2020-11-15 22:30:00+01:00	46400.0	11.84	11.96
2020-11-15 23:00:00+01:00	48600.0	11.75	11.94
2020-11-15 23:30:00+01:00	49400.0	11.64	11.92

	minuteofday	dayofweek	minuteofday_cos \
2020-11-01 00:00:00+01:00	0	6	1.000000
2020-11-01 00:30:00+01:00	30	6	0.991445
2020-11-01 01:00:00+01:00	60	6	0.965926
2020-11-01 01:30:00+01:00	90	6	0.923880
2020-11-01 02:00:00+01:00	120	6	0.866025
...	...	...	...
2020-11-15 21:30:00+01:00	1290	6	0.793353
2020-11-15 22:00:00+01:00	1320	6	0.866025
2020-11-15 22:30:00+01:00	1350	6	0.923880
2020-11-15 23:00:00+01:00	1380	6	0.965926
2020-11-15 23:30:00+01:00	1410	6	0.991445

	minuteofday_sin	dayofweek_cos	dayofweek_sin \
2020-11-01 00:00:00+01:00	0.000000	0.62349	-0.781831
2020-11-01 00:30:00+01:00	0.130526	0.62349	-0.781831
2020-11-01 01:00:00+01:00	0.258819	0.62349	-0.781831
2020-11-01 01:30:00+01:00	0.382683	0.62349	-0.781831
2020-11-01 02:00:00+01:00	0.500000	0.62349	-0.781831

...	...	...	...
2020-11-15 21:30:00+01:00	-0.608761	0.62349	-0.781831
2020-11-15 22:00:00+01:00	-0.500000	0.62349	-0.781831
2020-11-15 22:30:00+01:00	-0.382683	0.62349	-0.781831
2020-11-15 23:00:00+01:00	-0.258819	0.62349	-0.781831
2020-11-15 23:30:00+01:00	-0.130526	0.62349	-0.781831

	dayofyear_cos	dayofyear_sin	lockdown	\
2020-11-01 00:00:00+01:00	0.500000	-0.866025	0.0	
2020-11-01 00:30:00+01:00	0.500000	-0.866025	0.0	
2020-11-01 01:00:00+01:00	0.500000	-0.866025	0.0	
2020-11-01 01:30:00+01:00	0.500000	-0.866025	0.0	
2020-11-01 02:00:00+01:00	0.500000	-0.866025	0.0	

...	...	...	...
2020-11-15 21:30:00+01:00	0.691771	-0.722117	0.0
2020-11-15 22:00:00+01:00	0.691771	-0.722117	0.0
2020-11-15 22:30:00+01:00	0.691771	-0.722117	0.0
2020-11-15 23:00:00+01:00	0.691771	-0.722117	0.0
2020-11-15 23:30:00+01:00	0.691771	-0.722117	0.0

	public_holiday	nb_school_areas_off	\
2020-11-01 00:00:00+01:00	0.0	0.0	
2020-11-01 00:30:00+01:00	0.0	0.0	
2020-11-01 01:00:00+01:00	0.0	0.0	
2020-11-01 01:30:00+01:00	0.0	0.0	
2020-11-01 02:00:00+01:00	0.0	0.0	

...	...	...
2020-11-15 21:30:00+01:00	0.0	0.0
2020-11-15 22:00:00+01:00	0.0	0.0
2020-11-15 22:30:00+01:00	0.0	0.0
2020-11-15 23:00:00+01:00	0.0	0.0
2020-11-15 23:30:00+01:00	0.0	0.0



	extra_long_weekend
2020-11-01 00:00:00+01:00	0.0
2020-11-01 00:30:00+01:00	0.0
2020-11-01 01:00:00+01:00	0.0
2020-11-01 01:30:00+01:00	0.0
2020-11-01 02:00:00+01:00	0.0

...	...
2020-11-15 21:30:00+01:00	0.0
2020-11-15 22:00:00+01:00	0.0
2020-11-15 22:30:00+01:00	0.0
2020-11-15 23:00:00+01:00	0.0
2020-11-15 23:30:00+01:00	0.0

[720 rows x 17 columns]

```
[31]: # compare portfolio forecast to reality
for c in ["contracts_count", "kva"]:
    to_plot = pd.merge(
        portfolio[(portfolio.index >= '2020-09-01') & (portfolio.index <=
        ↪ '2020-11-16')][c].to_frame(c+"_historic"),
        benchmark_test[c].to_frame(c+"_forecast"),
        how='outer', left_index=True, right_index=True
    )

    to_plot.plot(figsize=(16, 4))
```



Lets define some algorithms then train and predict with them. All the models we define implement the `enda.estimators.EndaEstimator` abstract class (see the docs).

Enda comes with wrappers around scikit-learn and H2O estimators :

- sklearn: `enda.ml_backends.sklearn_estimator.EndaSklearnEstimator`
- H2O: `enda.ml_backends.h2o_estimator.EndaH2OEstimator`

```
[32]: import time
import h2o
import random
import numpy
```

```

from sklearn.linear_model import LinearRegression, SGDRegressor
from sklearn.neural_network import MLPRegressor
from sklearn.ensemble import AdaBoostRegressor
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler

from enda.ml_backends.h2o_estimator import EndaH2OEstimator # enda's wrapper
    ↪ around H2O models
from h2o.estimators import H2OGeneralizedLinearEstimator
from h2o.estimators import H2OXGBoostEstimator
from h2o.estimators import H2OGradientBoostingEstimator
from h2o.estimators import H2ORandomForestEstimator
from h2o.estimators import H2ODeepLearningEstimator

```

```

[33]: random.seed(17) # set random seed for reproducibility
      numpy.random.seed(17) # for sklearn
      # for h2o we will define it in each model

```

```

[34]: all_models = dict()

```

```

[35]: # Some models with the sklearn machine learning backend

all_models['sklearn_lin_reg'] = EndaSklearnEstimator(LinearRegression())

all_models['sklearn_sgd'] = EndaSklearnEstimator(
    Pipeline([('standard_scaler', StandardScaler()),
              ('sgd', SGDRegressor())
            ])
)

all_models['sklearn_ada_boost'] = EndaSklearnEstimator(AdaBoostRegressor(
    n_estimators=500,
    loss='square',
    learning_rate=0.8)
)

all_models['sklearn_nn'] = EndaSklearnEstimator(
    Pipeline([('standard_scaler', StandardScaler()),
              ('mlp', MLPRegressor(
                  solver='adam',
                  activation='relu',
                  hidden_layer_sizes=[48, 48, 24],
                  max_iter=150
              ))
            ])
)

```

```
)
```

[36]: *# Some models with the h2o machine learning backend*

```
all_models['h2o_glm'] = EndaH2OEstimator(H2OGeneralizedLinearEstimator(
    standardize=False,
    intercept=True,
    seed=17)
)

all_models['h2o_rf'] = EndaH2OEstimator(H2ORandomForestEstimator(
    ntrees=300,
    max_depth=15,
    sample_rate=0.8,
    min_rows=10,
    nbins=52,
    mtries=3,
    seed=17
))

all_models['h2o_gbm'] = EndaH2OEstimator(H2OGradientBoostingEstimator(
    ntrees=500,
    max_depth=5,
    sample_rate=0.5,
    min_rows=5,
    seed=17
))

all_models['h2o_nn'] = EndaH2OEstimator(H2ODeepLearningEstimator(
    **{
        "activation": "Tanh",
        "hidden": [48, 48, 24],
        "distribution": "gaussian",
        "epochs": 20,
        "seed": 17
    })
)
```

[37]: *# You can add more models to the benchmark here if you like*

[38]: *# to train or predict with H2O models, we boot up a local h2o server*  
h2o.init(nthreads=-1)  
h2o.no\_progress()

Checking whether there is an H2O instance running at http://localhost:54321...  
not found.

Attempting to start a local H2O server...

Java Version: openjdk version "21.0.1" 2023-10-17 LTS; OpenJDK Runtime



```

Environment Zulu21.30+15-CA (build 21.0.1+12-LTS); OpenJDK 64-Bit Server VM
Zulu21.30+15-CA (build 21.0.1+12-LTS, mixed mode, sharing)
Starting server from /Users/clement.jeannesson/.pyenv/versions/3.9.10/envs/end
a_1.0.0_dev/lib/python3.9/site-packages/h2o/backend/bin/h2o.jar
Ice root: /var/folders/pp/kyc80_js50g283hj0_c4yrhc0000gp/T/tmp9w493lg0
JVM stdout: /var/folders/pp/kyc80_js50g283hj0_c4yrhc0000gp/T/tmp9w493lg0/h2o_c
lement_jeannesson_started_from_python.out
JVM stderr: /var/folders/pp/kyc80_js50g283hj0_c4yrhc0000gp/T/tmp9w493lg0/h2o_c
lement_jeannesson_started_from_python.err
Server is running at http://127.0.0.1:54321
Connecting to H2O server at http://127.0.0.1:54321 ... successful.

```

```

-----
H2O_cluster_uptime:      02 secs
H2O_cluster_timezone:    Europe/Paris
H2O_data_parsing_timezone: UTC
H2O_cluster_version:     3.46.0.1
H2O_cluster_version_age: 13 days
H2O_cluster_name:        H2O_from_python_clement_jeannesson_xqjnib
H2O_cluster_total_nodes: 1
H2O_cluster_free_memory: 3.984 Gb
H2O_cluster_total_cores: 8
H2O_cluster_allowed_cores: 8
H2O_cluster_status:      locked, healthy
H2O_connection_url:       http://127.0.0.1:54321
H2O_connection_proxy:     {"http": null, "https": null}
H2O_internal_security:    False
Python_version:           3.9.10 final
-----

```

```
[39]: benchmark_train.iloc[0:10, 2:5]
```

```
[39]:
```

	load_kw	tso_forecast_load_mw	t_weighted
2015-01-01 00:00:00+01:00	2490.925806	72900.0	-0.41
2015-01-01 00:30:00+01:00	2412.623113	71600.0	-0.48
2015-01-01 01:00:00+01:00	2365.611276	69900.0	-0.55
2015-01-01 01:30:00+01:00	2336.141065	70600.0	-0.66
2015-01-01 02:00:00+01:00	2300.935642	70500.0	-0.78
2015-01-01 02:30:00+01:00	2226.613719	69000.0	-0.89
2015-01-01 03:00:00+01:00	2166.173069	67200.0	-1.00
2015-01-01 03:30:00+01:00	2104.404493	65400.0	-1.11
2015-01-01 04:00:00+01:00	2064.678631	63800.0	-1.22
2015-01-01 04:30:00+01:00	2035.268532	62700.0	-1.25

```
[40]: benchmark_train.iloc[0:10, 2:5].to_csv("training_test.csv")
```

```
[41]: # this should take between 5 and 15 minutes to run (in function of your
↳ hardware)
```

```

print("Benchmark with {} models : {}\n".format(len(all_models), list(all_models.
↪keys()))))
for model_name, model in all_models.items():
    model_start_time = time.time()
    print("Training {} before predicting with it..".format(model_name))
    model.train(benchmark_train, target_col='load_kw')
    model_prediction = model.predict(benchmark_test, target_col='load_kw')
    benchmark[model_name] = model_prediction
    print("{} took {:.1f} seconds.\n".format(model_name, time.
↪time()-model_start_time))

```

Benchmark with 8 models : ['sklearn\_lin\_reg', 'sklearn\_sgd',  
'sklearn\_ada\_boost', 'sklearn\_nn', 'h2o\_glm', 'h2o\_rf', 'h2o\_gbm', 'h2o\_nn']

Training sklearn\_lin\_reg before predicting with it..  
sklearn\_lin\_reg took 0.1 seconds.

Training sklearn\_sgd before predicting with it..  
sklearn\_sgd took 0.9 seconds.

Training sklearn\_ada\_boost before predicting with it..  
sklearn\_ada\_boost took 44.4 seconds.

Training sklearn\_nn before predicting with it..  
sklearn\_nn took 91.2 seconds.

Training h2o\_glm before predicting with it..

<IPython.core.display.HTML object>

h2o\_glm took 3.4 seconds.

Training h2o\_rf before predicting with it..  
h2o\_rf took 15.6 seconds.

Training h2o\_gbm before predicting with it..  
h2o\_gbm took 8.2 seconds.

Training h2o\_nn before predicting with it..  
h2o\_nn took 18.3 seconds.





[42]: benchmark





[42]:		actual_load_kw	sklearn_lin_reg	sklearn_sgd	\
	2020-11-01 00:00:00+01:00	6817.332090	7116.262400	7268.504320	
	2020-11-01 00:30:00+01:00	6326.667322	6896.504843	7046.769174	
	2020-11-01 01:00:00+01:00	6172.223671	6682.516424	6830.605028	

2020-11-01 01:30:00+01:00	6050.575318	6699.648917	6847.511206
2020-11-01 02:00:00+01:00	5898.881230	6635.737998	6782.325729
...	...	...	...
2020-11-15 21:30:00+01:00	7657.293444	7647.649943	7784.966759
2020-11-15 22:00:00+01:00	7317.540759	7516.196417	7653.352240
2020-11-15 22:30:00+01:00	7580.051439	7599.955734	7738.868558
2020-11-15 23:00:00+01:00	7496.273993	7784.720105	7926.057651
2020-11-15 23:30:00+01:00	7376.005701	7838.739518	7981.093363
	sklearn_ada_boost	sklearn_nn	h2o_glm \
2020-11-01 00:00:00+01:00	7004.055238	6978.222468	7411.859182
2020-11-01 00:30:00+01:00	6800.756289	6668.168002	7173.717434
2020-11-01 01:00:00+01:00	6457.625390	6356.734892	6935.575686
2020-11-01 01:30:00+01:00	6462.474690	6243.347042	6958.317705
2020-11-01 02:00:00+01:00	6430.078251	6060.438533	6878.974772
...	...	...	...
2020-11-15 21:30:00+01:00	7723.935769	8126.484696	7259.411106
2020-11-15 22:00:00+01:00	7688.075642	7852.186456	7146.039856
2020-11-15 22:30:00+01:00	7751.777186	7770.791269	7282.209600
2020-11-15 23:00:00+01:00	7841.638268	7786.286328	7531.807069
2020-11-15 23:30:00+01:00	7846.176657	7719.989282	7622.605723
	h2o_rf	h2o_gbm	h2o_nn
2020-11-01 00:00:00+01:00	6618.429014	7238.106976	7134.318405
2020-11-01 00:30:00+01:00	6244.503840	6695.938679	6792.191181
2020-11-01 01:00:00+01:00	6056.103537	6378.550857	6490.857203
2020-11-01 01:30:00+01:00	5971.874814	6277.017389	6340.915400
2020-11-01 02:00:00+01:00	5830.723221	6196.947266	6155.570844
...	...	...	...
2020-11-15 21:30:00+01:00	7495.115877	7419.043766	8127.375499
2020-11-15 22:00:00+01:00	7393.640487	7141.754804	7934.052204
2020-11-15 22:30:00+01:00	7428.028846	7417.634841	7856.795770
2020-11-15 23:00:00+01:00	7434.956921	7220.599850	7806.171075
2020-11-15 23:30:00+01:00	7442.448993	7264.606470	7709.001933

[720 rows x 9 columns]

```
[43]: # visualize predictions
for c in benchmark.columns:
    if c != "actual_load_kw":
        to_plot = benchmark[["actual_load_kw", c]]
        to_plot.plot(ylim=0, figsize=(16, 4))
```





```
[44]: # compute the mean absolute percentage error of each algo
scoring = enda.Scoring(predictions_df=benchmark, target="actual_load_kw")
scoring.mean_absolute_percentage_error().to_frame("mape")
```

```
[44]:
```

	mape
sklearn_lin_reg	6.972668
sklearn_sgd	7.448227
sklearn_ada_boost	6.687777
sklearn_nn	4.821875
h2o_glm	9.056221
h2o_rf	2.832520
h2o_gbm	2.155135
h2o_nn	5.559592

## 1.6 6. Benchmark with Backtesting

In traditional machine learning, we need more than just 1 evaluation to test an algorithm. We typically use cross-validation to see if the algorithm is not biased and if it can be expected to work well in most cases. For time-series predictions we cannot do a regular cross-validation because it is not realistic : we always want to train using historical data that happened before the prediction.

Here we will do **backtesting** week after week. With the given dataset, this means : - for each week w from early 2019 until the end of the dataset : train using data from the beginning of the dataset (early 2015) until a few days before week w, then eval on w. - the first iteration will train an algorithm using data from 2015 to 2018, then eval on the first week of 2019 - the second iteration will train using data from 2015 to a bit before the first week of 2019, then eval on the second week of 2019 - and so on... - keep the predictions of each time-step using this method, from early 2019 to november 2020.

- then compare these predictions to the historic data to evaluate the quality of each algorithm.

This makes most sense if in your production environment, you plan to retrain the algorithm regularly with recent data.

Backtesting can take a significant amount of time. We backtest only 2 linear regressions below in order to have an example that runs fast. Don't hesitate to add other algorithms.

```
[45]: all_models = dict()

all_models['sklearn_lin_reg'] = EndaSklearnEstimator(LinearRegression())

all_models['h2o_glm'] = _
    ↪EndaH2OEstimator(H2OGeneralizedLinearEstimator(standardize=False, _
    ↪intercept=True))
```

```
[46]: from dateutil.relativedelta import relativedelta
portfolio_train_length = relativedelta(months=1)
```

```

[47]: start_backtesting_dt = pd.to_datetime('2019-01-01 00:00:00+01:00').
      ↪tz_convert('Europe/Paris')
benchmark = historic[historic.index>=start_backtesting_dt]["load_kw"].
      ↪to_frame("actual_load_kw")
days_in_each_iteration = 28

for model_name, model in all_models.items():

    count_iterations = 0
    model_predictions = []
    for train_set, test_set in enda.BackTesting.yield_train_test(
        historic,
        start_eval_datetime=start_backtesting_dt,
        days_between_trains=days_in_each_iteration,
        gap_days_between_train_and_eval=14
    ):
        count_iterations += 1
        if count_iterations <= 2 or count_iterations % 10 == 0:
            print("Model {}, backtesting iteration {}, train set {}->{}, test_
            ↪set {}->{}\n".format(
                model_name, count_iterations,
                train_set.index.min(), train_set.index.max(),
                test_set.index.min(), test_set.index.max()))

        # featurize
        test_set = test_set.drop(columns=["load_kw"])

        # forecast portfolio for the test_set
        pf_train_start = enda.TimezoneUtils.add_interval_to_day_dt(
            day_dt=test_set.index.min(),
            interval=-portfolio_train_length,
        )
        pf_train = portfolio[(portfolio.index >= pf_train_start) & (portfolio.
        ↪index < test_set.index.min())]

        forecast_portfolio = enda.Contracts.forecast_portfolio_linear(
            portfolio_df=pf_train,
            start_forecast_date=test_set.index.min(),
            end_forecast_date_exclusive=test_set.index.
            ↪max()+relativedelta(minutes=30),
            freq='30min',
            tzinfo='Europe/Paris'
        ) # recent portfolio trend

        test_set['kva'] = forecast_portfolio['kva']
        test_set['contracts_count'] = forecast_portfolio['contracts_count']

```

```

# train and predict
model.train(train_set, target_col='load_kw')
model_predictions.append(model.predict(test_set, target_col='load_kw'))

benchmark[model_name] = pd.concat(model_predictions)

```

```

Model sklearn_lin_reg, backtesting iteration 1, train set 2015-01-01
00:00:00+01:00->2018-12-17 23:30:00+01:00, test set 2019-01-01
00:00:00+01:00->2019-01-28 23:30:00+01:00

```

```

Model sklearn_lin_reg, backtesting iteration 2, train set 2015-01-01
00:00:00+01:00->2019-01-14 23:30:00+01:00, test set 2019-01-29
00:00:00+01:00->2019-02-25 23:30:00+01:00

```

```

Model sklearn_lin_reg, backtesting iteration 10, train set 2015-01-01
00:00:00+01:00->2019-08-26 23:30:00+02:00, test set 2019-09-10
00:00:00+02:00->2019-10-07 23:30:00+02:00

```

```

Model sklearn_lin_reg, backtesting iteration 20, train set 2015-01-01
00:00:00+01:00->2020-06-01 23:30:00+02:00, test set 2020-06-16
00:00:00+02:00->2020-07-13 23:30:00+02:00

```

```

Model h2o_glm, backtesting iteration 1, train set 2015-01-01
00:00:00+01:00->2018-12-17 23:30:00+01:00, test set 2019-01-01
00:00:00+01:00->2019-01-28 23:30:00+01:00

```

```

Model h2o_glm, backtesting iteration 2, train set 2015-01-01
00:00:00+01:00->2019-01-14 23:30:00+01:00, test set 2019-01-29
00:00:00+01:00->2019-02-25 23:30:00+01:00

```

```

Model h2o_glm, backtesting iteration 10, train set 2015-01-01
00:00:00+01:00->2019-08-26 23:30:00+02:00, test set 2019-09-10
00:00:00+02:00->2019-10-07 23:30:00+02:00

```

```

Model h2o_glm, backtesting iteration 20, train set 2015-01-01
00:00:00+01:00->2020-06-01 23:30:00+02:00, test set 2020-06-16
00:00:00+02:00->2020-07-13 23:30:00+02:00



```

```

[48]: # visualize predictions
for c in benchmark.columns:
    if c != "actual_load_kw":
        to_plot = benchmark[["actual_load_kw", c]]
        to_plot.plot(ylim=0, figsize=(16, 4))

```





```
[49]: # compute mean absolute percentage error
scoring = enda.Scoring(predictions_df=benchmark, target="actual_load_kw")
scoring.mean_absolute_percentage_error().to_frame("mape")
```

```
[49]:
```

	mape
sklearn_lin_reg	13.830003
h2o_glm	14.362408

If you have time/computing power: - try more algorithms in the backtesting benchmark, this is longer but more reliable than a simple benchmark (think of it as crossval versus single eval in a non-time-series setup). - reduce the “days\_in\_each\_iteration” down to 7 if you think you can have a weekly training in your production environment.

## 1.7 7. Make the prediction

Seeing the results from just the basic benchmark, we here decide to predict using h2o’s gbm (and our set of hyperparameters). We now need to train it on the full dataset and make the prediction.

In the input data, the TSO forecast is only available for the next 7 days but the weather forecast is available for the next 11 days.

We use **EndaEstimatorWithFallback** to be able to predict with or without TSO data.

Checkout more EndaEstimators here: <https://github.com/enercoop/enda/blob/main/enda/estimators.py>. They work on top of all supported machine learning backends.

```
[50]: from enda.estimators import EndaEstimatorWithFallback
```

```
[51]: # create the forecast_input_data dataframe

# we will forecast the portfolio for the next 11 days
forecast_portfolio = enda.Contracts.forecast_portfolio_linear(
    portfolio_df=portfolio[portfolio.index >= '2020-11-01 00:00:00+01:00'],
    start_forecast_date=pd.to_datetime("2020-12-01 00:00:00+01:00").
    ↪tz_convert("Europe/Paris"),
    end_forecast_date_exclusive=pd.to_datetime("2020-12-12 00:00:00+01:00").
    ↪tz_convert("Europe/Paris"),
    freq='30min',
    tzinfo='Europe/Paris'
)

# this time we don't remove rows where tso_forecast is missing
forecast_input_data = pd.merge(
    forecast_portfolio,
    weather_and_tso_forecasts,
    how='inner', left_index=True, right_index=True
)

# add feature engineering
forecast_input_data = featurize_datetime(forecast_input_data)
forecast_input_data
```

```
[51]:
```

	contracts_count	kva \	
2020-12-01 00:00:00+01:00	96046.000857	819322.806873	
2020-12-01 00:30:00+01:00	96046.650461	819329.517545	
2020-12-01 01:00:00+01:00	96047.300064	819336.228217	
2020-12-01 01:30:00+01:00	96047.949668	819342.938889	
2020-12-01 02:00:00+01:00	96048.599272	819349.649561	
...	...	...	
2020-12-11 21:30:00+01:00	96385.743543	822832.488295	
2020-12-11 22:00:00+01:00	96386.393147	822839.198967	
2020-12-11 22:30:00+01:00	96387.042751	822845.909639	
2020-12-11 23:00:00+01:00	96387.692354	822852.620311	
2020-12-11 23:30:00+01:00	96388.341958	822859.330983	

	tso_forecast_load_mw	t_weighted	t_smooth \
2020-12-01 00:00:00+01:00	66100.0	4.69	5.08
2020-12-01 00:30:00+01:00	64200.0	4.82	5.10
2020-12-01 01:00:00+01:00	61900.0	4.96	5.12
2020-12-01 01:30:00+01:00	62800.0	5.04	5.13
2020-12-01 02:00:00+01:00	62300.0	5.13	5.14
...	...	...	...
2020-12-11 21:30:00+01:00	NaN	8.25	6.03
2020-12-11 22:00:00+01:00	NaN	8.22	5.94

2020-12-11 22:30:00+01:00	NaN	8.16	5.83
2020-12-11 23:00:00+01:00	NaN	8.11	5.78
2020-12-11 23:30:00+01:00	NaN	8.11	5.73

	minuteofday	dayofweek	minuteofday_cos \
2020-12-01 00:00:00+01:00	0	1	1.000000
2020-12-01 00:30:00+01:00	30	1	0.991445
2020-12-01 01:00:00+01:00	60	1	0.965926
2020-12-01 01:30:00+01:00	90	1	0.923880
2020-12-01 02:00:00+01:00	120	1	0.866025
...	...	...	...
2020-12-11 21:30:00+01:00	1290	4	0.793353
2020-12-11 22:00:00+01:00	1320	4	0.866025
2020-12-11 22:30:00+01:00	1350	4	0.923880
2020-12-11 23:00:00+01:00	1380	4	0.965926
2020-12-11 23:30:00+01:00	1410	4	0.991445

	minuteofday_sin	dayofweek_cos	dayofweek_sin \
2020-12-01 00:00:00+01:00	0.000000	0.623490	0.781831
2020-12-01 00:30:00+01:00	0.130526	0.623490	0.781831
2020-12-01 01:00:00+01:00	0.258819	0.623490	0.781831
2020-12-01 01:30:00+01:00	0.382683	0.623490	0.781831
2020-12-01 02:00:00+01:00	0.500000	0.623490	0.781831
...	...	...	...
2020-12-11 21:30:00+01:00	-0.608761	-0.900969	-0.433884
2020-12-11 22:00:00+01:00	-0.500000	-0.900969	-0.433884
2020-12-11 22:30:00+01:00	-0.382683	-0.900969	-0.433884
2020-12-11 23:00:00+01:00	-0.258819	-0.900969	-0.433884
2020-12-11 23:30:00+01:00	-0.130526	-0.900969	-0.433884

	dayofyear_cos	dayofyear_sin	lockdown \
2020-12-01 00:00:00+01:00	0.861702	-0.507415	0.0
2020-12-01 00:30:00+01:00	0.861702	-0.507415	0.0
2020-12-01 01:00:00+01:00	0.861702	-0.507415	0.0
2020-12-01 01:30:00+01:00	0.861702	-0.507415	0.0
2020-12-01 02:00:00+01:00	0.861702	-0.507415	0.0
...	...	...	...
2020-12-11 21:30:00+01:00	0.935717	-0.352752	0.0
2020-12-11 22:00:00+01:00	0.935717	-0.352752	0.0
2020-12-11 22:30:00+01:00	0.935717	-0.352752	0.0
2020-12-11 23:00:00+01:00	0.935717	-0.352752	0.0
2020-12-11 23:30:00+01:00	0.935717	-0.352752	0.0

	public_holiday	nb_school_areas_off \
2020-12-01 00:00:00+01:00	0.0	0.0
2020-12-01 00:30:00+01:00	0.0	0.0
2020-12-01 01:00:00+01:00	0.0	0.0


2020-12-01 01:30:00+01:00	0.0	0.0
2020-12-01 02:00:00+01:00	0.0	0.0
...	...	...
2020-12-11 21:30:00+01:00	0.0	0.0
2020-12-11 22:00:00+01:00	0.0	0.0
2020-12-11 22:30:00+01:00	0.0	0.0
2020-12-11 23:00:00+01:00	0.0	0.0
2020-12-11 23:30:00+01:00	0.0	0.0


	extra_long_weekend
2020-12-01 00:00:00+01:00	0.0
2020-12-01 00:30:00+01:00	0.0
2020-12-01 01:00:00+01:00	0.0
2020-12-01 01:30:00+01:00	0.0
2020-12-01 02:00:00+01:00	0.0
...	...
2020-12-11 21:30:00+01:00	0.0
2020-12-11 22:00:00+01:00	0.0
2020-12-11 22:30:00+01:00	0.0
2020-12-11 23:00:00+01:00	0.0
2020-12-11 23:30:00+01:00	0.0

[528 rows x 17 columns]

```
[52]: # show recent portfolio and forecast
for c in ["contracts_count", "kva"]:
    to_plot = pd.merge(
        portfolio[(portfolio.index >= '2020-10-01')][c].to_frame(c+"_historic"),
        forecast_input_data[c].to_frame(c+"_forecast"),
        how='outer', left_index=True, right_index=True
    )

    to_plot.plot(figsize=(16, 4))
```





```
[53]: # tso data is missing after 2020-12-07 :
forecast_input_data[forecast_input_data.index>='2020-12-07 23:00:00+01:00'].
      ↪head()
```

```
[53]:
```

	contracts_count	kva \
2020-12-07 23:00:00+01:00	96262.968462	821564.171299
2020-12-07 23:30:00+01:00	96263.618065	821570.881971
2020-12-08 00:00:00+01:00	96264.267669	821577.592643
2020-12-08 00:30:00+01:00	96264.917273	821584.303315
2020-12-08 01:00:00+01:00	96265.566876	821591.013987

	tso_forecast_load_mw	t_weighted	t_smooth \
2020-12-07 23:00:00+01:00	70200.0	3.94	4.07
2020-12-07 23:30:00+01:00	69600.0	3.94	4.07
2020-12-08 00:00:00+01:00	NaN	3.95	4.07
2020-12-08 00:30:00+01:00	NaN	3.88	4.06
2020-12-08 01:00:00+01:00	NaN	3.81	4.05

	minuteofday	dayofweek	minuteofday_cos \
2020-12-07 23:00:00+01:00	1380	0	0.965926
2020-12-07 23:30:00+01:00	1410	0	0.991445
2020-12-08 00:00:00+01:00	0	1	1.000000
2020-12-08 00:30:00+01:00	30	1	0.991445
2020-12-08 01:00:00+01:00	60	1	0.965926

	minuteofday_sin	dayofweek_cos	dayofweek_sin \
2020-12-07 23:00:00+01:00	-0.258819	1.00000	0.000000
2020-12-07 23:30:00+01:00	-0.130526	1.00000	0.000000
2020-12-08 00:00:00+01:00	0.000000	0.62349	0.781831
2020-12-08 00:30:00+01:00	0.130526	0.62349	0.781831
2020-12-08 01:00:00+01:00	0.258819	0.62349	0.781831

	dayofyear_cos	dayofyear_sin	lockdown \
2020-12-07 23:00:00+01:00	0.909308	-0.416125	0.0

2020-12-07 23:30:00+01:00	0.909308	-0.416125	0.0
2020-12-08 00:00:00+01:00	0.916317	-0.400454	0.0
2020-12-08 00:30:00+01:00	0.916317	-0.400454	0.0
2020-12-08 01:00:00+01:00	0.916317	-0.400454	0.0

	public_holiday	nb_school_areas_off	\
2020-12-07 23:00:00+01:00	0.0	0.0	
2020-12-07 23:30:00+01:00	0.0	0.0	
2020-12-08 00:00:00+01:00	0.0	0.0	
2020-12-08 00:30:00+01:00	0.0	0.0	
2020-12-08 01:00:00+01:00	0.0	0.0	

	extra_long_weekend
2020-12-07 23:00:00+01:00	0.0
2020-12-07 23:30:00+01:00	0.0
2020-12-08 00:00:00+01:00	0.0
2020-12-08 00:30:00+01:00	0.0
2020-12-08 01:00:00+01:00	0.0

```
[54]: gbm_1 = EndaH2OEstimator(H2OGradientBoostingEstimator(
    ntrees=500,
    max_depth=5,
    sample_rate=0.5,
    min_rows=5
))

gbm_2 = EndaH2OEstimator(H2OGradientBoostingEstimator(
    ntrees=500,
    max_depth=5,
    sample_rate=0.5,
    min_rows=5
))

m = EndaEstimatorWithFallback(
    resilient_column="tso_forecast_load_mw",
    estimator_with=gbm_1,
    estimator_without=gbm_2
)
```

```
[55]: m.train(full_train_set, target_col='load_kw')
```

```
[56]: import joblib
model_file_path = os.path.join(DIR, "gbm_with_fallback.pickle")
```

```
[57]: # save the model for later
joblib.dump(m, filename=model_file_path)
```

```
[57]: ['./gbm_with_fallback.pickle']
```

```
[58]: del m
```

```
[59]: # load the model from disk (works even if you shutdown then restarted the H2O_
      ↪server)
      m2 = joblib.load(filename=model_file_path)
```

```
[60]: m_prediction = m2.predict(forecast_input_data, target_col="load_kw")
```

```
[61]: # a good prediction is made until 2020-12-11
      # even where TSO forecast is missing
      m_prediction.tail()
```


```
[61]:
```

	load_kw
2020-12-11 21:30:00+01:00	9787.209992
2020-12-11 22:00:00+01:00	9375.710344
2020-12-11 22:30:00+01:00	9541.818818
2020-12-11 23:00:00+01:00	9403.112938
2020-12-11 23:30:00+01:00	9208.277543

```
[ ]:
```

```
[62]: # visualize recent load along with our forecast; remember we don't have recent_
      ↪actual load so there is a time-gap.
      # (remember that the prediction takes weather forecast and more information_
      ↪into account)
      to_plot = pd.merge(
          historic["load_kw"][-1000:].to_frame("historic_kw"),
          m_prediction.rename(columns={"load_kw": "forecast_kw"}),
          how='outer', left_index=True, right_index=True
      )
      to_plot.plot(ylim=0, figsize=(16, 4))
```

```
[62]: <Axes: >
```



```
[63]: # don't forget to shutdown your h2o local server  
      h2o.cluster().shutdown()  
      # wait for h2o to really finish shutting down  
      time.sleep(5)
```

H2O session \_sid\_9c51 closed.

## 1.8 Conclusion

That's all for Example B. Check out Example C next. Thanks for reading and don't hesitate to send feedback at: [team-data@enercoop.org](mailto:team-data@enercoop.org) !